

PennOS

Project 2 Group 11

<https://github.com/CIS548/22fa-project-2-group-11>

Reifon Chiu, Sumukh Govindaraju, Jio Jeong, Rowena Lu

File Structure

```
bin/
  pennfat
  pennos
log/
  scheduler.log
doc/
  CompanionDoc.pdf
src/
  fat/
    fat_util.c
    fat_util.h
    fat.c
    file_kernel_funcs.c
    file_kernel_funcs.h
  kernel/
    os.c
    pcb_table.c
    pcb_table.h
    process_kernel_funcs.c
    process_kernel_funcs.h
    scheduler.c
    scheduler.h
    threads.c
    threads.h
  lib/
    directory_entry.c
    directory_entry.h
    errno.c
    errno.h
    fd.c
    fd.h
    file_system.h
    linked_list.c
```

```
linked_list.h
log.c
log.h
macros.h
parser.h
pcb.c
pcb.h
signals.h
status.c
status.h
shell/
  commands.c
  commands.h
  job.c
  job.h
  redirects.c
  redirects.h
  shell.c
  shell.h
user/
  file_user_funcs.c
  file_user_funcs.h
  process_user_funcs.c
  process_user_funcs.h
  scheduler_user_funcs.c
  scheduler_user_funcs.h
  stress.c
  stress.h
.gitignore
Makefile
parser-aarch64.o
parser-x86_64.o
README.md
```

System Calls

Process Related System Calls

pid_t p_spawn(void (*func) (), char *argv[], int fd0, int fd1)

This function calls k_process_create and returns the pid of the created child_process.

If the child_process is NULL, then the return value is -1. Additionally, fd0 and fd1 are

the file descriptors of the input and output file respectively. This can be found in `process_user_funcs.c`.

pid_t p_waitpid(pid_t pid, int *wstatus, bool nohang)

If the given pid is -1, wait on any child that changes status. Otherwise, wait on the child with the given pid. If there isn't a child process for the parent to wait on, then the return value is -1. If nohang is set to true, and if any child or the specified child has changed status, return the pid. Otherwise, return 0. If nohang is set to false, then the current process is blocked until the child changes state, at which point the pid is eventually returned. This function also has various logging events depending on the child's status and whether nohang is false. This can be found in `process_user_funcs.c`.

int p_kill(pid_t pid, int sig)

This function calls `k_process_kill`, and sends the given signal to the process with the given pid. Returns 0 upon success. If the process with the given pid is not found, return -1. This can be found in `process_user_funcs.c`.

void p_exit(void)

This calls `k_process_kill` with a `S_SIGTERM` and exits the current thread unconditionally. If the active job is null, the function returns sets the error number and returns. This can be found in `process_user_funcs.c`.

int p_nice(pid_t pid, int priority)

This sets the priority of the thread pid to priority. This returns 0 if there are no errors. This returns -1 when the given priority is not -1, 0, or 1 and when there is no process with the given pid in the process table. This can be found in `scheduler_user_funcs.c`.

void p_sleep(unsigned int ticks)

This sets the calling process to be blocked for ticks ticks, and then sets the thread to running. If the active job is null, the function returns sets the error number and returns. This can be found in `scheduler_user_funcs.c`.

File System System Calls

int f_open(const char *fname, int mode)

This opens a file with the name fname in the mode mode, and returns a file descriptor of that file. The mode can be any one of three modes: `F_WRITE` or 0, `F_READ` or 1, and `F_APPEND` or 2. For `F_WRITE`, the file is opened for reading and writing and is overwritten, and returns an error if the same file is being opened more than once. For `F_READ`, the file is opened only for reading, and returns an error if the file does not exist. For `F_APPEND`, the file is opened for reading and writing but does not overwrite the existing data. Also note that for `F_APPEND` the pointer is at the end of the file. Additionally, `F_WRITE` and `F_APPEND` creates the file if it does not exist. Lastly, this returns a file descriptor upon

success, but a negative value in the case that there's an error (invalid characters in filenames, file is not able to be created, etc.).

int f_read(int fd, int n, char *buf)

This returns the number of n bytes referenced by the file descriptor fd. If EOF is reached, this returns 0. If there's an error with reading the file (ex: file is not opened / not found), this returns -1.

int f_write(int fd, const char *str, int n)

This returns the number of n bytes referenced by the string str that is written to the file descriptor fd, and increments the file pointer by n. This returns the number of bytes written, or -1 in the case of an error (ex: file is not opened / not found, or there is no more space).

int f_close(int fd)

This closes the file fd and returns 0 upon success, but -1 upon failure.

int f_unlink(const char *fname)

This removes the file fname and returns 0 upon success, but -1 if the file has not been closed.

int f_lseek(int fd, int offset, int whence)

This sets the file pointer for fd to be the offset relative to whence. Additionally, whence can be any of the constants F_SEEK_SET, F_SEEK_CUR, and F_SEEK_END, which are the beginning of the file, the current position of the file pointer, and the end of file respectively. This returns 0 upon success, and -1 in the case of an error (ex: file is not opened / not found).

int f_ls(const char *filename)

This lists the file filename in the directory, but lists all the files in the current directory if filename is NULL. This returns 1 if the filename is found, otherwise it returns 0.

PCB Description

The PCB contains several fields:

Current status of the job.

process_status status

Whether the status has changed or not since last calling waitpid() on this process.

bool status_changed

Whether the process is blocked or not

bool blocked

Whether the process is in the background or not

bool is_bg

Context of the process.

ucontext_t* context

PID for the process.

pid_t pid

Parent process's PCB.

struct pcb_st* parent

Dynamically allocated string representing the pipeline name.

char* cmd

Linked list of zombie child processes.

linked_list* zombieLL

Linked list of nonzombie child processes.

linked_list* childLL

Priority level.

int priority

Array of open file descriptors with fd[0] = STDIN and fd[1] = STDOUT.

int fd[1024]

Dynamically allocated and null terminated array of string arguments passed to the process.

char argv**

Note also that the PCB's process_status can be any of the following: STOPPED, RUNNING, ZOMBIED, TERMINATED, EXITED, and ORPHANED (please see status.h for more details).

Process Table Description

The process table is a linked list that contains pcb data structures of all processes (see pcb_table.c)

The processes associated with the process table is the following:

void init_pcb_table()

This function initializes the linked list for the process table.

void add_pcb_to_table(pcb* to_add)

This function adds the specified pcb into the process table.

pcb* find_pcb_in_table(int pid)

This function finds the pcb of the process specified by the input pid within the process table.

void delete_pcb_in_table(int pid)

This function deletes the pcb of the process specified by the input pid within the process table.

linked_list* get_process_table()

This function returns the address of the process table.

void free_process_table()

This function frees all the dynamically allocated variables (pcb and fields of pcb) associated with the process table.

PCB Related Functions

In `pcb.c`, there are the following functions:

bool pid_equal_predicate(void* pid, void *target_pcb);

This is used as a predicate in a `linked_list` function that searches for the `linked_list_elem` with the given pid in the given pcb.

pcb* k_process_create(pcb* parent)

This creates the child thread along with its PCB. Additionally, it retains the properties of the parent, and returns a reference to the new PCB. In the case of a `malloc` error, `NULL` is returned and an error message is printed. This can be found in `process_kernel_funcs.c`.

void k_process_kill(pcb *process, int signal)

This sends the signal `signal` to the process, and logs the event. If the signal is `S_SIGTERM`, it zombies the target process and it orphans all of its children by iterating through its children linked lists and setting their status as `ORPHANED`). This can be found in `process_kernel_funcs.c`.

void k_process_cleanup(pcb* process)

This function iterates through the initial process 's children and if they're in the zombie linked list or in the non-zombie linked list, the children are removed from the lists and their PCB 's are freed. In the case that the children's status are neither TERMINATED, EXITED, nor ORPHANED, an error is printed stating that the process is not a zombie. This function helps to clean the memory that's allocated from `k_process_create()`. This returns nothing. This can be found in `process_kernel_funcs.c`.

void free_process_pcb(pcb *process)

This function goes through a process pcb frees each of its values as long as they aren't NULL.

bool process_complete(pcb* j)

This function returns whether the pid of the pcb j is equal to FINISHED_PID_VAL.

The remaining functions have to deal with the linked list which is used for the job queues.

pcb* get_job_from_pid(linked_list *linked_list, int pid)

This function returns the process pcb given a pid.

job* get_current_job(linked_list *linked_list)

This function returns the last stopped background job, or the last background job if nothing is stopped. The background job queue is passed in as a linked list.

void update_job_status(linked_list *linked_list, int pid, int status)

This function changes the status of a job to the given status , given the pid of the job.

Signal Description

The signals we used were the default signals as described in the writeup. These signals are: S_SIGSTOP, S_SIGCONT, and S_SIGTERM.

Error Numbers

```
// Process/kernel errors
NOERROR, // No error
NOCHILDCREATED, // k_process_create failed to create a shell
ACTIVEJOBNULL, // Passed in job is null
NOTINPCBTABLE, // Job is invalid and not in PCB Table.
KILLZOMBIE, // Tried to send signal to non-existent/zombie process.
NOSUCHCHILD, // Parent doesn't have child with PID.
NOTFOUNDINSCHEDULER, // Job could not be found in the scheduler with the
right priority
INVALIDPRIORITY, // Invalid priority integer
INVALIDSIGNAL, // Invalid signal integer passed to p_kill

// File kernel errors
INVALID_WHENCE, // Invalid whence value for f_lseek
INVALID_OFFSET, // Invalid offset for f_lseek
FILE_NOT_FOUND, // File is not found in file user function
UNALLOCATED_BLOCK, // File does not have an allocated first block
PERMISSION_DENIED, // User does not have permissions to access the file

// File user errors
INVALID_MODE, // Invalid integer file mode
INVALID_FILE_NAME, // File name is NULL
INVALID_FILE_NAME_POSIX, // File name does not follow POSIX standards
ATTEMPTED_DOUBLE_WRITE, // Trying to open again in write mode
READ_FILE_NOT_FOUND, // f_open File not found and was opened in F_READ
mode
CLOSE_UNOPEN_FILE, // f_close File is not open
DOUBLE_DELETION, // File has already been deleted
FILE_NOT_FOUND_OFT, // fd not found in OFT

// Critical Errors
NO_MORE_SPACE // no more space in the fs
```


Description of Remaining Files

fat/

fat_util.c, fat_util.h, fat.c, file_kernel_funcs.c, file_kernel_funcs.h

kernel/

os.c, pcb_table.c, pcb_table.h, process_kernel_funcs.c, process_kernel_funcs.h, scheduler.c, scheduler.h, threads.c, threads.h

lib/

directory_entry.c, directory_entry.h, errno.c, errno.h, fd.c, fd.h, file_system.h, log.c, log.h, macros.h, parser.h, signals.h, status.c, status.h

shell/

commands.c, commands.h, job.c, job.h, redirects.c, redirects.h, shell.c, shell.h

user/

scheduler_user_funcs.c, scheduler_user_funcs.h, stress.c, stress.h

Description of Other Data Structures

Directory Entry

This is a struct that we created that contains data pertaining to each directory entry. The fields are as follows:

name: the name of the directory entry

size: the size of the file

firstBlock: the first block of the file

type: the type of the file casted from the FilyType enum

perm: the file's read/write/execute permissions

mtime: the file's most recent modification time

reserved: 16 reserved bytes

Open File Table

We use an **Open File Table** linked list to keep track of which files are open. Each element is simply a File Descriptor, as described below.

File Descriptor

This is a struct that we used that contains the data pertaining to each file descriptor, which is used in a linked list as our OFT, as described above. The fields are as follows:

***de:** a pointer to the corresponding directory entry

ind: the integer (fd) assigned to this file descriptor

ref_index: the number of fd array elements that point to this node

mode: the mode (F_READ, F_WRITE, or F_APPEND)
f_pos: the file offset from the start of the file with respect to the file system
d_pos: the file offset of the directory entry of the file being opened with respect to the file system

Jobs

This is a struct that we created to keep track of the background jobs. The fields are as follows:

job_pid: the pid of the process
job_id: Job index to print in jobs built in command

The related functions are the following:

job* create_job(pid_t pid): creates a job object using the given pid
void free_job(void* j): frees the dynamically allocated pointers in the specified job.
bool job_equal_predicate(void *target_pid, void *curr_job): used as a predicate in a linked_list function that searches for the linked_list_elem with the given pid and the given job
bool job_id_equal_predicate(void *target_job_id, void* curr_job): used as a predicate in a linked_list function that searches for the linked_list_elem with the given job_id and the given job

Linked List

This is a fairly standard generalized linked list. This can be found in linked_list.c, along with the header file linked_list.h .

Extra Credit

Memory Leaks

There should be no memory leaks in the standalone FAT or Penn OS.